

# CS 521: Trustworthy AI Systems

## MP1

Adharsh Kamath Raghupathi (netID: ak128)

### Problem 1:

#### Part (1)

Given the original input  $x$ , the target class  $t$ , and the perturbation magnitude  $\epsilon$ , we can compute FGSM adversary as follows:

```
adv_x = x - eps * x.grad.sign()
print("Adversarial Example: ", adv_x)
```

With this snippet (in the file <https://github.com/adharshkamath/cs521-sp25/blob/main/fgsm.py>), we can see the following output:

```
Original Input:  tensor([[0.6584, 0.8175, 0.1262, 0.0541, 0.2268, 0.8405, 0.5393, 0.9798, 0.8597,
                          0.5977]], requires_grad=True)
Original Class:  2
Adversarial Example:  tensor([[ 0.1584,  1.3175,  0.6262, -0.4459,  0.7268,  0.3405,  0.0393,  0.4798,
                          1.3597,  0.0977]], grad_fn=<SubBackward0>)
Adversarial Class:  0
tensor(0.5000)
```

As we can see, the FGSM attack found an adversarial example that is close enough to the original input. but is classified into a different class.

#### Part (2)

The original value of epsilon (0.5) did not work for the FGSM attack, for target class 1. I tried different values of epsilon but that did not help. So I tried the Iterative FGSM attack. But I had to increase the epsilon to 0.85, and use  $\alpha = 0.25$  with max iterations set to 10. The relevant code snippet is as follows:

```
epsReal = 0.85
eps = epsReal - 1e-7
alpha = 0.25
iterations = 10

for iteration in range(iterations):
    adv_x.requires_grad_(True)
    N.zero_grad()
    loss = L(N(adv_x), torch.tensor([t], dtype=torch.long))
    loss.backward()
    with torch.no_grad():
        adv_x = adv_x - (alpha * adv_x.grad.sign())
        delta = adv_x - x
        delta = torch.clamp(delta, -eps, eps) # clamp to eps ball
        adv_x = x + delta
```

```
new_class = N(adv_x).argmax(dim=1).item()
distance = torch.norm((x-adv_x), p=float('inf')).data
print(f"Iteration {iteration}: Distance {distance}")
if new_class == t:
    print("Attack successful")
    break
```

The rest of the code can be found in the script here:  
<https://github.com/adharshkamath/cs521-sp25/blob/main/fgsm.py>

## Problem 2:

### Part (1)

Thy Google colab notebook with the code can be found [here](#).

Below are the standard and robust accuracies of the three models ( $L_\infty$  attack with  $\text{eps}=8/255$  and  $L_2$  attack with  $\text{eps} = 0.75$ )

- pretr\_Linf.pth: Standard= 82.80%,  $L_2$  robustness=45.96%,  $L_\infty$  robustness = 50.65%
- pretr\_L2.pth: Standard=88.75%,  $L_2$  robustness=53.25%,  $L_\infty$  robustness = 29.15%,
- pretr\_RAMP.pth: Standard=81.19%,  $L_2$  robustness=59.25%,  $L_\infty$  robustness = 48.93%,

From the numbers we can see that  $L_2$  robustness training does not hold up well against the  $L_\infty$  PGD attack. The model with  $L_\infty$  robustness training holds up well against both  $L_2$  and  $L_\infty$  attacks. The model with RAMP training holds up well against both the attacks but has a slightly lower standard accuracy compared to the other two.

### Part (2)

The code for this is in the same Google colab notebook linked above. Below are the accuracies of the of the networks for multi-norm robustness:

- pretr\_Linf.pth: 47.05%
- pretr\_L2.pth: 30.85%
- pretr\_RAMP.pth: 49.76%

From the accuracies we can see that RAMP trained model performs best, followed by the model that is  $L_\infty$  robustness trained. Finally  $L_2$  robustness trained model does not perform as great as either of the two. So, RAMP training helps the most in this case.

### Problem 3:

The paper introduces a range of new unforeseen adversaries, including 18 new non- $L_P$  attacks. They illustrate the point that it is not always possible to anticipate the full range of attacks that a model might be subjected to in the real world. They propose a new framework, ImageNet-UA, that aims to evaluate models against unforeseen adversaries. They show that existing models are lacking, when evaluated using this framework.